

AI value-drift as a portfolio risk

A one-page brief for a private equity operating partner. Where AI value-drift shows up across a portfolio, a portfolio risk lens, seven questions worth asking a portfolio CEO or CRO, and a paragraph ready to drop into a value-creation plan. Designed to read in five minutes before a monthly portfolio review.



Cumai Aboul Housn, Iga Sloan and Alex Abbott

Supero · Published 2 May 2026 · Updated 14 July 2026 · 5 min read

AI is now embedded inside the revenue engine of most companies in a portfolio. It is shaping lead scores, forecasts, pricing, and customer prioritisation, quietly and asset by asset. For an operating partner, the question is not whether each management team is using AI well. It is whether the portfolio has the visibility, the standard, and the questions in hand to protect the value-creation plan, and the exit multiple, before the drift shows up in the numbers.

KEY TERMS

Value-drift in the revenue engine

The slow shift in what "good" looks like inside a company because AI-mediated workflows are steering behaviour before leadership has consciously decided what to reward. Across a portfolio it compounds asset by asset, unseen on the dashboards.

AI governance inventory

A register of the AI systems in use across a company, each with its decision influence, data dependency, EU AI Act risk classification, named human owner, and reversal authority. The core diligence artefact for AI, at entry and at exit.

Exit-value exposure

The discount a buyer or lender applies when a company cannot evidence how AI shapes its commercial decisions. The gap between a dashboard that looks healthy and a revenue engine that can be explained under diligence.

Reversal authority

The named individual with the standing to suspend, retrain, or withdraw an AI deployment that underperforms or drifts. A core test of governance maturity, and one of the cheapest to check across a portfolio.

Portfolio-ready paragraph

A short, defensible statement of position on AI value-drift, written so an operating partner can lift it directly into a value-creation plan, a board pack, or a monthly portfolio review without further drafting.

1 The position in three lines

The risk is not the absence of AI in the portfolio. It is the absence of a shared standard for AI that has already arrived inside every asset. Each company has bought AI separately, configured it locally, and, in most cases, owns it nowhere with a full view of consequence. What looks like productivity on a monthly dashboard can be value-drift, the revenue engine slowly optimising for something other than the strategy the value-creation plan is underwriting.

For a PE-backed company, or for an operating partner supporting a portfolio, this is not only an operational issue. It is also a board issue, a diligence issue, and eventually an **exit-value issue**. The same AI inventory that reduces regulatory exposure today becomes evidence of control, maturity, and explainability in future investor or buyer diligence. The operating partner is one of the few people with the standing to set that standard once and apply it across every asset.

2 The value at risk, not just the breach cost

With operational improvement now driving more of the return than leverage and multiple expansion, the first exposure is not a breach. It is upside quietly not landing: AI inside the revenue engine optimising for something other than the thesis the value-creation plan underwrites. The downside numbers are real too. Both belong on the same page.

200–500 bps

MARGIN, PLUS 5–10% REVENUE GROWTH, A DISCIPLINED AI PROGRAMME CAN ADD TO A PORTFOLIO COMPANY OVER THE HOLD, IF IT IS POINTED AT THE PLAN RATHER THAN AT WHATEVER CONVERTS FASTEST.

7.8 months

TIME TO IMPACT FOR PRICING, THE FASTEST VALUE LEVER, AT A 4% FAILURE RATE. VALUE-DRIFT IS THE TAX ON THAT SPEED WHEN AI IS SHAPING PRICE, MIX AND PRIORITISATION OFF-THESIS.

63%

SHARE OF ORGANISATIONS WITH NO POLICY GOVERNING AI USAGE, EXPOSED BY DEFAULT, AND WITH NO LINE OF SIGHT TO WHETHER AI IS HELPING OR DRIFTING (IBM, 2025).

€620k

AVERAGE BREACH PREMIUM CARRIED BY ORGANISATIONS WITH HIGH SHADOW-AI USE (ROUGHLY \$670K, IBM, 2025), ON TOP OF >50% OF AI-ATTRIBUTED LAYOFFS FORRESTER EXPECTS TO BE REVERSED.

The upside lines never reach a board pack because no one is measuring drift; the downside lines only reach it once an incident forces them to. The operating partner's advantage is that the same short set of questions surfaces both, at the same cadence, across every company in the fund.

Sources: margin and revenue uplift, Gorilla Logic AI value-creation analysis, 2026; pricing lever speed and failure rate, Simon-Kucher PE Value Creation Study, 2025; policy and breach figures, IBM Cost of a Data Breach, 2025; layoff reversal, Forrester. Operational improvement now drives the majority of PE value creation (Bain; Apollo; McKinsey Global Private Markets Report, 2026).

3 The portfolio risk lens

Three lenses, each with a worst-case an operating partner should be prepared to discuss at the investment committee, not just at the portfolio board.

Value creation

Value-drift in the revenue engine. AI shaping lead scores, forecasts, pricing and customer prioritisation in ways that drift from the ICP and strategy the value-creation plan is underwriting.

A company on plan by activity and dashboard, but off plan on win rate, deal size, and retention, drifting hardest in the final year of the hold, when the largest share of exit-year EBITDA is meant to be made.

Diligence & exit

The ability to evidence, on demand, how AI influences commercial decisions, who owns each system, and how it is classified and controlled.

A buyer or lender diligence team finds AI shaping revenue with no inventory, no owners, and no documentation, and prices the uncertainty as a discount to the multiple.

Regulatory & reputational

EU AI Act risk classification and literacy obligations across the portfolio; AI decisions affecting customers, candidates or employees with no defensible rationale.

A high-risk system identified post-deployment in one asset, documentation reconstructed retrospectively, in front of a regulator, an LP, or a press cycle that has already asked.

4 Seven questions to ask a portfolio CEO or CRO

None of these questions are hostile. They protect the management team and the value-creation plan. The first asks where AI is meant to be creating value; the rest make sure it still is. Ask them the same way in every company, and the portfolio becomes comparable.

1 Where is AI supposed to be adding growth or margin in this company, is that tied to a specific line in the value-creation plan, and can you show me it is actually landing, not just running?

2 Where is AI already making, or materially shaping, commercial decisions in your revenue engine, and can you give me that as a list with a named human owner for each, not a discussion?

- 3 When did you last check that all the AI-powered use cases supporting your GTM operating model are still optimising for the ICP and the strategy in the value-creation plan, rather than for whatever converts fastest?
- 4 If you needed to pause or withdraw any AI system from the revenue engine tomorrow, how quickly could you do it, and who has the authority to sign that off?
- 5 Which of your AI systems sit in the EU AI Act risk classification, what oversight does that require, and are you ahead of the deadlines or behind them?
- 6 If a buyer's diligence team asked you to evidence how AI shapes your commercial decisions, what could you produce in 48 hours, and what would you rather they did not find?
- 7 What is your written policy on AI usage at work, and how do you know your people are operating inside it rather than improvising?

5 Value-creation-plan-ready paragraph

If you want to record a portfolio-wide position, the following paragraph is drafted to be transferred directly into a value-creation plan, a portfolio board pack, or a monthly review.

SUGGESTED WORDING

AI is now embedded across the productivity and revenue-generating decisions of the portfolio, from the tools that shape the commercial pipeline to those that manage people and operations. As a standard condition of the value-creation plan, each company will maintain and present, at the agreed cadence, a register of the AI use cases in operation, each ranked by its level of risk under the EU AI Act and mapped to the date enforcement takes effect, with a named human owner accountable for each and evidence of meaningful human oversight. Management teams will confirm that AI literacy obligations, in force since 2 February 2025, are met; that they are ready for the general application of the Act from 2 August 2026; and that any high-risk use case is on a plan to meet the 2 December 2027 deadline (2 August 2028 for AI embedded in regulated products). A process for monitoring value-drift against the approved commercial strategy will be agreed and assigned. Any material AI use case without documented ownership, risk classification, or oversight will be flagged with a remediation timeline, so that the same inventory that manages regulatory and operational risk today also stands as diligence-ready evidence of control and maturity at exit.

Apply the same standard across every asset. The substance, inventory, classification, ownership, oversight, literacy, drift, remediation, is what turns a scattered set of AI pilots into a defensible position at the investment committee and in the data room.

The conversation this brief is built for

The most expensive AI decisions across a portfolio in the next twelve months will be the ones no operating partner ever saw, already inside an approved system, configured by a function the board rarely meets, owned by no one with authority to intervene. The remedy is not more reporting from each asset. It is one standard, asked for the same way across every company, so drift surfaces while it is still cheap to correct, and the AI inventory becomes an exit proof point rather than a diligence discount.

This is the standard Supero sets. Alex Abbott runs the diagnostic with each management team and holds the portfolio-wide line. Cumai Aboul Housn carries the technical authority for the AI inventory and the infrastructure beneath it. Iga Sloan brings the EU AI Act risk classification and the documentation that survives buyer and lender diligence, and turns the regulatory clock from a liability into a plan. Together they are the standard, and the people who install it.

What that looks like in practice: one AI value-drift diagnostic, run the same way across every asset in the fund, producing a ranked red-amber-green scorecard for the investment committee and a value-creation-plan clause for each portfolio company. The fund sets the standard once; each company executes it. Fixed scope, comparable output, done before the next quarterly review.

About the authors



Cumai Aboul Housn

Enterprise Transformation Architect, AI and Digital Innovation

Cumai carries the technical authority for the AI inventory that sits behind a value-creation plan. He helps portfolio companies translate AI use into structured, EU AI Act aligned, diligence-ready evidence of control.



Iga Sloan

EU AI Act and Regulatory Strategy

Iga brings the regulatory acumen that helps operating partners treat the EU AI Act as a springboard for performance rather than a compliance cost. She classifies each AI use case by risk, maps it to the enforcement clock, and turns balanced governance into a competitive advantage the value-creation plan can underwrite and a buyer can price.



Alex Abbott

Founder and Chief Revenue Officer, Supero

Alex runs the diagnostic with management teams and works with operating partners to set one AI oversight standard across a portfolio, so revenue-engine decisions stay documented, defensible, and on plan through to exit.